

# Teaching reproducibility: motivation, direct instruction and practice

*Joel Ostblom      Tiffany A. Timbers*

## Introduction

The data science definition that we embrace in the master of data science program at UBC and in the undergrad data science courses that we're developing there is the study and development of reproducible and audible processes to obtain insight from data. When you go into the data science classroom students are usually very excited about learning data science but what they're most excited about is the second part of the definition, the insight from data part.

Often they are not even aware about the reproducible and audible processes part and they see that more as a pain/inconvenience. So you have this barrier when you're teaching the reproducibility aspects of data science. This probably arises because they likely do not even know what reproducibility is, and even if they do know about it, it is not the thing that is obviously/directly providing insight and so they're not excited about it.

Then we have this third challenge, which is that the tools that we use for reproducibility are not necessarily smooth and easy to learn. They usually have a pretty steep learning curve. Over our five years of teaching these things at UBC we've found some key things that we've experienced at least for teaching reproducibility successfully:

1. placing extra emphasis on motivation
2. direct instruction
3. lots of practice

In this paper, we will discuss why we believe each of these are important, give some high-level examples of how we do these, and then we give one detailed example of how we do each in our courses.

## Placing extra emphasis on motivation

Why do we need extra motivation when teaching reproducibility, compared to some other data science topics, such as machine learning? We think this is because students do not have intrinsic excitement or motivation on the topic of

reproducibility, they have little prior knowledge on this topic, and reproducibility concepts and in particular tools are challenging to learn.

For a specific example, the version control software that's the most commonly used one for reproducibility, Git, is notorious for being difficult to learn (Figure 1). Furthermore, there are many anecdotes that most people don't actually learn it deeply and they just get by trying a variety of commands until they find some things that work. Which can lead to users getting themselves into a lot of trouble. Sometimes this trouble is so difficult to get out of, that some professional data scientists and data science educators, for example Jenny Bryan, recommend the practice of "burning it all down" and starting from scratch - which really defeats many of the purposes of version control.



Fig. 1: Infamous xkcd comic that highlights the difficulty of learning and using the version control software Git.

um our markdown is a lovely tool but  
 anytime you involve lottic anything  
 into anything um you can get some very  
 challenging error messages  
 that aren't really clear about where  
 they come from and what you should do  
 next so this is  
 another place that students find  
 problems  
 and then some another tool that we use  
 in reproducibility has containerization

with docker  
and this is a really challenging one to  
teach and motivate students to learn  
because  
it's so different from writing code to  
do analysis because it takes a long time  
to install things  
and it takes a long time to automate the  
process of installing things and making  
a compute environment  
accessible and it's not that exciting  
because things already work on their own  
computer  
so hopefully these cartoons have  
illustrated um  
basically a lot of the motivation  
challenges we have with these tools  
so what do we do to motivate um in our  
classrooms  
so uh one thing that we've used is  
telling stories from the trenches  
so um those of us who teach these  
courses  
uh usually have had some experience with  
doing research and their phds or postdoc  
or are still currently doing research  
um and through these are lived  
experiences of learning reproducibility  
tools and applying it to our research we  
have made mistakes  
our collaborators have made mistakes and

we can share these with our students  
in the master of data science program a  
lot of the  
students have work experience that may  
have touched on to data before  
and so they also have stories from the  
trenches and so uh in the  
you know in person classroom you can do  
some think pair share exercises around  
this  
and get the students to talk about their  
stories as well as hear your stories  
um and then in the zoom classroom you  
can always do breakout rooms for this  
sort of thing  
another example that was something that  
i learned from jenny bryan  
is let them fail but let them do this in  
a controlled manner  
so um at least myself i experienced a  
lot of failure in graduate school in my  
postdoctoral research in reproducibility  
and uh it took a long time and it really  
slowed me down  
and i'll argue in a few slides later  
that we shouldn't do this  
um so if we can set up these scenarios  
where they feel a little bit of this  
pain  
but it's for a short period of time that  
can be very useful

this can be especially useful when  
you're teaching undergraduates data  
science because they haven't had  
these real life experiences  
and then third is something that i'm  
going to try and bring into the new  
course that i'm developing  
which are case studies of failures uh  
that have had real world consequences  
um and again these are you know the  
master students  
this hasn't been as is important for i  
don't think because  
they've had some of these real life  
experiences themselves at work  
um but for undergraduate students again  
that haven't had the opportunity to do  
research  
um don't really understand what the real  
world consequences are so i think  
real case studies um that have real  
world consequences  
are a good way to to motivate for that  
so here's an example of uh in the  
classroom i let it fail in a controlled  
manner so  
i've done this exercise in the course  
that i teach that's called workflows for  
data science  
in this course um it's a project-based  
course but we have to teach them all the

skills so they can do a reproducible data analysis and at the end of the course one of the things that we get them to do is to make their compute environments for their analysis reproducible using a tool called docker docker has this is not an easy thing to teach or learn it's a lot of overhead you're teaching that we need to teach them to write docker files you know writing how you install things in linux and half of the students are windows users so like there is a lot of you know barrier and stickiness to teaching this subject so you really need to motivate it so one way that i do this in a demonstration is that we give students a data analysis project pipeline on github and we asked them go to go to that github repository read the instructions and try and replicate the analysis our students already have some git skills at this point in the course so they're able to do the git clone they're reading the readme trying to follow along

uh with the instructions to run it um in  
in the first instance  
uh we've given them an analysis and  
we've like intentionally put in a lot of  
packages that we know they don't have  
installed  
like we've got off into  
into crayon or into pipey eye and like  
found some very bizarre packages we  
don't even necessarily use them but you  
just tuck them in there  
um and if you want to be really you know  
uh  
sneaky you tuck them in there like in  
the middle of the script so they're not  
even at the top of the scripts  
and you ask them to to work on getting  
that analysis to run  
and it takes some time and eventually  
they figure it out but it's  
it's frustrating and then you give them  
the same analysis  
in a different github repository um but  
you give that it has a  
it has a docker solution so there's a  
docker image that exists on docker hub  
um the readme gives clear instructions  
on how to run it and replicate the  
analysis  
and they're able to do this and they're  
able to replicate the analysis in a

couple minutes  
so this is an example of how you can a  
let them fail in a controlled manner  
and then at the same time within the  
same learning  
time period give them a solution and  
motivate them to that solution and then  
they're in a good mind frame that  
even when learning soccer beca is hard  
and challenging  
they have the motivation to learn it  
because that they know that their  
analysis is going to be more useful for  
other people  
afterwards

### **Direct instruction**

okay so the second thing that i said was  
important is direct instruction  
so why is direct instruction important  
well from um  
i think for those of us who've been  
using reproducibility  
tools in our research what i'm going to  
say here is probably  
not new to you but reproducibility  
is not something that most people or  
students figure out through exploration  
and  
inquiry um based learning or if you do  
it's not an efficient



way of doing it um there's a lot of  
you know we're using a lot of borrowed  
tools from software engineering that are  
being repurposed  
for um for for science and  
reproducibility  
and so a lot of the uh you know getting  
up and getting started has a lot of  
assumed knowledge behind it  
and um there's not a lot of like clear  
and easy on roads to  
these things and because um  
yeah i think and there's also  
how would i say this uh just because  
it's still fairly new i would say  
that there's there's not um  
a lot of culture around it that's like  
that's very common uh to like show  
people where these obvious armor apps  
are  
so i and then there's the challenge of  
the tools that i talked about in the  
previous example so i think again  
having some direct instruction is is  
important and i really love this excerpt  
from roger peng's blog post  
from a couple of years ago that he wrote  
on the theory of data analysis  
and he writes here that there is no need  
for a new data analyst to learn  
about reproducibility from experience we

don't need to lead a junior data  
analysis  
down a months long winding path of  
non-reproducible analysis  
until they are finally bitten by the  
non-reproducibility  
bug and therefore learn their lesson we  
can just tell them  
hey in the past we found it useful to  
make our data analysis reproducible  
here's a workflow guide for you to use  
in your own analysis  
within that one statement we can  
compress with over 20 years of  
experience  
we i think owe it to our students to  
directly instruct them with like  
the best practices that you know that  
the reproducibility community has  
arrived on  
to date and then show them how to use  
these tools explicitly  
so how do we use direct exam instruction  
in the classroom  
um so we do a lot of live demos so in  
the programming classes we do  
you know live coding to show how to use  
r and python um but then when we're  
talking about other tools like docker  
uh our markdown or jupyter for doing  
reproducible reports

using version control with git and github we do a lot of live demos with those two tools as well and in doing that it makes it i think uh obvious to this more obvious to the students of like how to use these things number one and number two um you make mistakes and that humanizes the experience of of working with these tools that are somewhat challenging because students make mistakes too and they see that the experts are also making mistakes and then you're able to usually make those mistakes usually make your mistakes where things are a little bit more difficult or are a little bit more sticky and it gives you more time to spend on that area of the topic and explain why you made the mistake and where the misconceptions come in and and fix them the other thing that we use are guided worksheets and tutorials so there's a lot us giving the live demos which is useful but you can't you know for uh all the time be up in front of people and it's good for people to um work through and

actively engage with material out  
themselves and get a little bit of uh  
practice but not you know enough totally  
like free-for-all do your own thing  
whatever you think is best but in a  
guided way so um we have a lot of those  
in the program  
um but one thing i want to say is be  
careful when you uh  
just there are some dragons when you  
when you teach this stuff  
so um because we're teaching things that  
involve graphical user interfaces  
because we teach things that are coming  
from  
software engineering it's a very  
fast-moving field and which means that  
every time i teach this stuff  
i need to go through it before i send it  
out to the students because something  
has changed and something or something  
has broken  
and i have to come up with a work around  
so a story from this past term  
is that github which is  
you know the largest code hosting  
repository in the world  
uh decided to i think very rightfully so  
changed the name of their default branch  
they switched it from master to main  
but that caused lots of things to break

that caused all of our notes to have to  
um be rewritten um and it's still like  
a half solved problem because we have a  
whole bunch of resources that are still  
sitting on master branch that we haven't  
been able to or will not be able to  
uh quickly change over so  
um it is really important to have this  
direct instruction  
but these are things that you're going  
to have to kind of do new in the  
classroom every year make a new demo  
every year  
curate your gut worksheets and your  
tutorials every year because otherwise  
they're going to quickly  
fall out of usefulness an example of  
direct instruction  
um for teaching version control so we  
teach version control in our very first  
year introduction to data science  
uh course um and so we do this in  
uh kind of like a three-pronged approach  
for directed instruction  
so we give them a textbook reading that  
they're  
able to to use this is something we have  
to update every year because the  
graphical user interface that we choose  
to use  
changes we don't teach the command line

for this in the first year because it's  
a bit too  
overwhelming i think for the students um  
then we do a live demonstration where  
they  
they watch us use the github website  
they watch us use the get gui they watch  
us move files and  
add and commit and push and pull and  
then finally they work through a guided  
worksheet  
that asks them to do the same thing that  
we just did  
and then ask them questions along the  
way uh to test that they  
like really understand like what is  
committing what is adding what is  
pulling what is pushing where is the  
work going  
and if folks are interested i have uh  
some links embedded in this talk  
that will take you to some of the  
examples or resources that i'm talking  
about here  
okay and so the final thing that i said  
um

### **Lot's of practice**

is lots of practice so why do we need  
lots of practice  
for reproducibility for learning

reproducibility workflows and  
tools well there are really two  
fundamental ways that we commit things  
to long-term memory  
one is one trial learning and that  
usually requires some sort of emotional  
impact so that's like sometimes it's  
traumatic events and sometimes it's  
really good positive emotional events  
that you had like a really  
great birthday or your wedding or  
something like that you don't need those  
things to be repeated multiple times so  
that you remember them  
but that's not most of the things in  
school most of the things in school  
we learn about are through this  
repetitive space training  
um and so uh the the you know the best  
way to commit  
uh something to long-term memory that's  
not really emotional  
is to revisit it and repeat it multiple  
times and have breaks between  
those things and so that lets you commit  
it to memory  
however i think you want to go even a  
step further with reproducibility  
because when we teach reproducibility  
workflows and skills as instructors  
we actually want students to do more

than just learn about these  
things we actually want them to use them  
and put them into practice  
um in the classroom outside of the  
classroom in their work after the  
classroom  
and so we actually want to change their  
habits or behaviors and it's it's quite  
i think  
important to realize that okay it's not  
just understanding an algorithm  
it's it's understanding the concepts  
behind something like version control  
understanding the concepts behind  
something like a shippable and shareable  
compute environment  
and then knowing how to use those things  
and then  
once you leave the classroom wanting and  
being able to use those things without  
like  
saying no that's too hard or too tricky  
you want them to just  
do it out of habit because that's what  
they usually do  
um so an aside uh just a little  
if people are interested a book that's  
really recently made me think about more  
how we can tangibly do this is called  
atomic habits by james clear  
he's done a really good job of like



bringing the science of  
habit building and behavior change uh  
into an accessible book  
and um i think that when we think about  
getting students practice and changing  
their behavior with  
reproducible skills and workflows  
there's a lot of really  
interesting insight from behavior change  
and psychology and habit building that  
we can that we can borrow  
so um now i'll talk more practically  
about so  
at least right now what are we doing in  
the classroom to embed  
lots of practice so  
what we do is when we do our live demos  
we don't just have  
us do it then we pause and say okay  
students your turn  
do what i just did and so they saw  
it and then they actually have to type  
it into the keyboard or click it their  
mouse around  
the graphical user interface so they  
they practice it that  
way then we have lots of low stakes  
assessments  
with small or short problems so um  
we've moved into a lot of flipped  
classroom

um in in at ubc or at least in the data science so our introduction to data science course is a primarily flipped classroom so um we have uh literate code documents that have uh automated tests in them that the students are answering all kinds of questions about the data science content and then they're very short little pieces that are well guided but they get immediate feedback um and these things aren't worth very much and they do a lot of them so they do two of them a week in the data science course in the master of data science um program we've also started implementing this in some of our classrooms and the students really like this practice and it helps them really prepare for things like larger assessments like quizzes and and their their lab homework but it gives them lots of practice and then the other thing that we do is the learning technologies and platforms that we use are built and use authentic data science reproducibility tools and so i'll give an example of that now so in almost all of the master beta

science courses so i'm talking about 20  
courses here  
21 credit courses  
so 21 month long courses we use  
version control particularly github as  
our course management system  
so the homework instructions and  
assignments are distributed to the  
students as github repositories  
and the only way that they can submit  
their homework is by putting their  
homework in that github repository  
so they have to go through the cloning  
procedure or at least be able to somehow  
download this from the  
github uh website and then they have to  
be able to hopefully through  
things like pushing and committing send  
their work back to github but they would  
at least have to interface with the  
github  
uh website to do this um  
to try and uh incentivize  
um the actual actually using the get  
machinery to interact with github  
we also put part of the marks of each  
of these assignments as to mechanics and  
so um  
we need to see for example like three  
commits associated with every single  
assignment because

um we think you know we're trying to build these good habits and practices around like there's reasons why you use version control not just to submit your homework but

to active as a backup or in case you want to go back in time and change things

so by the end of this program um the students have version controlled their work in over 80 different repositories um so they have a lot of so they're very practiced and very used to it and they're basically you know you want them to be able to do it in their sleep um almost and so that when they leave the program and they go to work somewhere else it's just natural it's just one of their habits at this point that if they're going to work on a project it's going to go under version control

so we do this using tools and here i've listened there's many tools now which is pretty cool we're not the only program doing this

at all there's many tools now for using github

as a classroom learning management system um and so i've listed a couple of them here

folks are interested so

## Conclusions

We think over the past five years from teaching in the Master of Data Science and our Introduction to Data Science course that key things for teaching reproducibility in the data science classroom are providing extra emphasis on motivation, providing direct instruction so it's not a mystery of how you get started and what you need to do, as well as lots and lots of practice so that we can not only teach them the material and the concepts but so that this actually changes their practices and their workflows and they will use it after leaving your classroom.

## Q & A's from talk

(leaving here for now in case there is anything inspirational to add to the paper)

so i'm happy to take questions now here

um

or uh you can tweet to me on twitter and

i'm happy to answer there and again i've

posted the link for the slides

thanks very much tiffany amy did you

want to share the question period or

shall i um i

can i'm just gonna check the thread

i don't see anything right now

anyone have anything to start off with

i have some if there's none um otherwise

maybe mina has

comments i'd be i'd love just to have

mina and tiffany just like

tell us everything you know between the

two of you

was there something that you started

doing that you i mean obviously

it's a evolution right the these these

all of these  
programs are just in the evolution stage  
uh was this uh  
something that you started doing that  
you've really moved away from  
yeah i would say the pandemics even  
placed a greater emphasis on this  
so when we started teaching the master  
of data science program  
um we had a small cohort and we were in  
person which allowed us to provide a lot  
of support  
but as you scale these things um  
having that intimate close support is  
more difficult and so in the very first  
year of the program  
um in the mbs program we have like very  
we we have this philosophy that they  
should be able to  
be somewhat experts of running stuff on  
their laptop we do teach them some cloud  
tools but  
you should be able to install your  
software stack and and be able to set up  
your path and these sorts of things  
and so um we have a pretty i'd say like  
intensive list of like 20 things that  
they need to install in the first week  
they need to use git to submit their  
homework in the first week  
and it's a bit overwhelming and it's a

lot and so  
um what we've kind of moved we've kind  
of like eased off on that  
and and moved um to we get them there  
but we take long  
we take longer now to doing that so uh  
we've set up this year a jupiter hub to  
have them work in a cloud-based setting  
for the first week or two  
and then after the first week or two is  
when we transition them to their own  
laptop  
so that we give time for the no like the  
the  
expectation that setting up everybody's  
system  
um is going to encounter some bugs and  
take some time and that's going to be  
tricky  
we've also for the first assignment we  
no longer ask them in the first  
assignment to  
submit to github um that's assignment  
two so that we have to give ourselves  
like a week or two  
to to get them up to speed for getting  
github  
um and so i think that uh that sort of  
thing  
has um has definitely changed and been  
inspired with

so first i started working in the master  
data science program and then i started  
teaching undergraduates  
and teaching undergrad graduates has  
made me have to reframe things and think  
about things differently  
um and think about like how do i remove  
barriers so that  
people you know maybe people who who are  
for whatever reason more sensitive to  
not feeling like data sciences for  
them i don't want them to drop my class  
because they couldn't install something  
um so i think yeah that's something  
that's changed a lot  
let's quickly get to meena's question  
where she needs to go um  
a question about have you seen any  
changes in computing experiences of  
students applying to your  
ms program um are more students coming  
in with familiarity with these tools or  
not yet  
yeah yeah that's a great question i do a  
little survey  
every year about like in in the first  
class  
i'm like what tools have you used before  
and usually about half of the students  
have used r  
maybe three quarters of the students



have used python almost all of them have  
seen jupiter  
almost none of them still have seen get  
in github um so it's really quite  
amazing  
that um  
computer science programming  
prerequisite uh so they have  
it's they don't have to have had our pro  
or python before  
um but yeah it's still interesting that  
even  
though i'll comment in a second that we  
are seeing people with more  
technical or data science skills coming  
to our program  
it's still the reproducibility  
experience with reproducibility related  
tools are  
are aside from jupiter like um not  
as present as one might expect i am  
seeing more and more  
uh people having like in data science  
applying to our program which is  
something kind of new and interesting  
for us to think about because our  
program was really designed  
not for somebody who is like a data  
science undergraduate like somebody who  
had  
an undergraduate in a different

discipline and wants to then apply  
you know data science to their  
discipline so we're still thinking about  
like how  
how we're going to handle the change of  
like there's going to be more and more  
undergraduates coming in with this  
expertise  
yeah it's a super interesting problem um  
john's asking  
um what to do with docker and windows  
um is there something special about  
documents yeah  
so it uh it can work  
um but everything with windows is a  
little bit more challenging  
uh so what my strategy is is i have  
um i i'm a mac person uh but i have a  
i also have a separate pc laptop where i  
have linux and windows installed  
so that before i teach every course i go  
through and make sure that  
i know how to install things on windows  
and  
what instructions to provide to students  
there's still always surprises  
um one thing we do this on quite a large  
scale largest scale with the master  
students but i'd say like  
we're dealing with 50 or 60 windows  
different windows laptops every

every year and so to make our lives  
easier  
um we've been really tightly restricting  
which version of windows  
that they have because uh then it's  
easier to know so we say you have to  
have windows 10  
you have to have this build um and it  
can't be windows home um basically it  
has to be enterprise pro  
or or education and by doing that that  
has reduced some problems but every year  
something new comes up like  
i can just tell you this week i'm  
teaching uh python packaging with poetry  
and git bash doesn't work with poetry  
anymore this year it worked last year  
but it doesn't work anymore there's a  
game  
that have issue open it's not resolved  
so now we're using anaconda prompt on  
the windows machines  
we have a solution um but it's  
it's just it's it's one of those there  
there will be dragons in this field  
yeah it's yeah keeping changing things  
it's just so much work right  
you think you're done and at least 20  
years ago right the folks they write out  
their theory equations and that was it  
they were done for the next 20 years

we've got to update ourselves every six  
months  
oh did you have any other closing  
comments or thoughts that you wanted to  
say