

Teaching Responsibility: motivation, direct instruction and practice

Joel Ostblom Tiffany A. Timbers

Introduction

The data science definition that we embrace in the master of data science program at UBC and in the undergrad data science courses that we're developing there is the study and development of reproducible and audible processes to obtain insight from data. When you go into the data science classroom students are usually very excited about learning data science but what they're most excited about is the second part of the definition, the insight from data part.

Often they are not even aware about the reproducible and audible processes part and they see that more as a pain/inconvenience. So you have this barrier when you're teaching the reproducibility aspects of data science. This probably arises because they likely do not even know what reproducibility is, and even if they do know about it, it is not the thing that is obviously/directly providing insight and so they're not excited about it.

Then we have this third challenge, which is that the tools that we use for reproducibility are not necessarily smooth and easy to learn. They usually have a pretty steep learning curve. Over our five years of teaching these things at UBC we've found some key things that we've experienced at least for teaching reproducibility successfully:

1. placing extra emphasis on motivation
2. direct instruction
3. lots of practice

In this paper, we will discuss why we believe each of these are important, give some high-level examples of how we do these, and then we give one detailed example of how we do each in our courses.

Placing extra emphasis on motivation

Why do we need extra motivation when teaching reproducibility, compared to some other data science topics, such as machine learning? We think this is because students do not have intrinsic excitement or motivation on the topic of

reproducibility, they have little prior knowledge on this topic, and reproducibility concepts and in particular tools are challenging to learn.

For a specific example, the version control software that's the most commonly used one for reproducibility, Git, is notorious for being difficult to learn (Figure 1). Furthermore, there are many anecdotes that most people don't actually learn it deeply and they just get by trying a variety of commands until they find some things that work. Which can lead to users getting themselves into a lot of trouble. Sometimes this trouble is so difficult to get out of, that some professional data scientists and data science educators, for example Jenny Bryan, recommend the practice of "burning it all down" and starting from scratch - which really defeats many of the purposes of version control.



Fig. 1: Infamous xkcd comic that highlights the difficulty of learning and using the version control software Git.

um our markdown is a lovely tool but
 anytime you involve lottic anything
 into anything um you can get some very
 challenging error messages
 that aren't really clear about where
 they come from and what you should do
 next so this is
 another place that students find
 problems
 and then some another tool that we use
 in reproducibility has containerization

with docker
and this is a really challenging one to
teach and motivate students to learn
because
it's so different from writing code to
do analysis because it takes a long time
to install things
and it takes a long time to automate the
process of installing things and making
a compute environment
accessible and it's not that exciting
because things already work on their own
computer
so hopefully these cartoons have
illustrated um
basically a lot of the motivation
challenges we have with these tools
so what do we do to motivate um in our
classrooms
so uh one thing that we've used is
telling stories from the trenches
so um those of us who teach these
courses
uh usually have had some experience with
doing research and their phds or postdoc
or are still currently doing research
um and through these are lived
experiences of learning reproducibility
tools and applying it to our research we
have made mistakes
our collaborators have made mistakes and

we can share these with our students
in the master of data science program a
lot of the
students have work experience that may
have touched on to data before
and so they also have stories from the
trenches and so uh in the
you know in person classroom you can do
some think pair share exercises around
this
and get the students to talk about their
stories as well as hear your stories
um and then in the zoom classroom you
can always do breakout rooms for this
sort of thing
another example that was something that
i learned from jenny bryan
is let them fail but let them do this in
a controlled manner
so um at least myself i experienced a
lot of failure in graduate school in my
postdoctoral research in reproducibility
and uh it took a long time and it really
slowed me down
and i'll argue in a few slides later
that we shouldn't do this
um so if we can set up these scenarios
where they feel a little bit of this
pain
but it's for a short period of time that
can be very useful

this can be especially useful when
you're teaching undergraduates data
science because they haven't had
these real life experiences
and then third is something that i'm
going to try and bring into the new
course that i'm developing
which are case studies of failures uh
that have had real world consequences
um and again these are you know the
master students
this hasn't been as is important for i
don't think because
they've had some of these real life
experiences themselves at work
um but for undergraduate students again
that haven't had the opportunity to do
research
um don't really understand what the real
world consequences are so i think
real case studies um that have real
world consequences
are a good way to to motivate for that
so here's an example of uh in the
classroom i let it fail in a controlled
manner so
i've done this exercise in the course
that i teach that's called workflows for
data science
in this course um it's a project-based
course but we have to teach them all the

skills so they can do a reproducible data analysis and at the end of the course one of the things that we get them to do is to make their compute environments for their analysis reproducible using a tool called docker docker has this is not an easy thing to teach or learn it's a lot of overhead you're teaching that we need to teach them to write docker files you know writing how you install things in linux and half of the students are windows users so like there is a lot of you know barrier and stickiness to teaching this subject so you really need to motivate it so one way that i do this in a demonstration is that we give students a data analysis project pipeline on github and we asked them go to go to that github repository read the instructions and try and replicate the analysis our students already have some git skills at this point in the course so they're able to do the git clone they're reading the readme trying to follow along

uh with the instructions to run it um in
in the first instance
uh we've given them an analysis and
we've like intentionally put in a lot of
packages that we know they don't have
installed
like we've got off into
into crayon or into pipey eye and like
found some very bizarre packages we
don't even necessarily use them but you
just tuck them in there
um and if you want to be really you know
uh
sneaky you tuck them in there like in
the middle of the script so they're not
even at the top of the scripts
and you ask them to to work on getting
that analysis to run
and it takes some time and eventually
they figure it out but it's
it's frustrating and then you give them
the same analysis
in a different github repository um but
you give that it has a
it has a docker solution so there's a
docker image that exists on docker hub
um the readme gives clear instructions
on how to run it and replicate the
analysis
and they're able to do this and they're
able to replicate the analysis in a

couple minutes
so this is an example of how you can a
let them fail in a controlled manner
and then at the same time within the
same learning
time period give them a solution and
motivate them to that solution and then
they're in a good mind frame that
even when learning soccer beca is hard
and challenging
they have the motivation to learn it
because that they know that their
analysis is going to be more useful for
other people
afterwards

Direct instruction

okay so the second thing that i said was
important is direct instruction
so why is direct instruction important
well from um
i think for those of us who've been
using reproducibility
tools in our research what i'm going to
say here is probably
not new to you but reproducibility
is not something that most people or
students figure out through exploration
and
inquiry um based learning or if you do
it's not an efficient

way of doing it um there's a lot of
you know we're using a lot of borrowed
tools from software engineering that are
being repurposed
for um for for science and
reproducibility
and so a lot of the uh you know getting
up and getting started has a lot of
assumed knowledge behind it
and um there's not a lot of like clear
and easy on roads to
these things and because um
yeah i think and there's also
how would i say this uh just because
it's still fairly new i would say
that there's there's not um
a lot of culture around it that's like
that's very common uh to like show
people where these obvious armor apps
are
so i and then there's the challenge of
the tools that i talked about in the
previous example so i think again
having some direct instruction is is
important and i really love this excerpt
from roger peng's blog post
from a couple of years ago that he wrote
on the theory of data analysis
and he writes here that there is no need
for a new data analyst to learn
about reproducibility from experience we

don't need to lead a junior data
analysis
down a months long winding path of
non-reproducible analysis
until they are finally bitten by the
non-reproducibility
bug and therefore learn their lesson we
can just tell them
hey in the past we found it useful to
make our data analysis reproducible
here's a workflow guide for you to use
in your own analysis
within that one statement we can
compress with over 20 years of
experience
we i think owe it to our students to
directly instruct them with like
the best practices that you know that
the reproducibility community has
arrived on
to date and then show them how to use
these tools explicitly
so how do we use direct exam instruction
in the classroom
um so we do a lot of live demos so in
the programming classes we do
you know live coding to show how to use
r and python um but then when we're
talking about other tools like docker
uh our markdown or jupyter for doing
reproducible reports

using version control with git and github we do a lot of live demos with those two tools as well and in doing that it makes it i think uh obvious to this more obvious to the students of like how to use these things number one and number two um you make mistakes and that humanizes the experience of of working with these tools that are somewhat challenging because students make mistakes too and they see that the experts are also making mistakes and then you're able to usually make those mistakes usually make your mistakes where things are a little bit more difficult or are a little bit more sticky and it gives you more time to spend on that area of the topic and explain why you made the mistake and where the misconceptions come in and and fix them the other thing that we use are guided worksheets and tutorials so there's a lot us giving the live demos which is useful but you can't you know for uh all the time be up in front of people and it's good for people to um work through and

actively engage with material out
themselves and get a little bit of uh
practice but not you know enough totally
like free-for-all do your own thing
whatever you think is best but in a
guided way so um we have a lot of those
in the program
um but one thing i want to say is be
careful when you uh
just there are some dragons when you
when you teach this stuff
so um because we're teaching things that
involve graphical user interfaces
because we teach things that are coming
from
software engineering it's a very
fast-moving field and which means that
every time i teach this stuff
i need to go through it before i send it
out to the students because something
has changed and something or something
has broken
and i have to come up with a work around
so a story from this past term
is that github which is
you know the largest code hosting
repository in the world
uh decided to i think very rightfully so
changed the name of their default branch
they switched it from master to main
but that caused lots of things to break

that caused all of our notes to have to
um be rewritten um and it's still like
a half solved problem because we have a
whole bunch of resources that are still
sitting on master branch that we haven't
been able to or will not be able to
uh quickly change over so
um it is really important to have this
direct instruction
but these are things that you're going
to have to kind of do new in the
classroom every year make a new demo
every year
curate your gut worksheets and your
tutorials every year because otherwise
they're going to quickly
fall out of usefulness an example of
direct instruction
um for teaching version control so we
teach version control in our very first
year introduction to data science
uh course um and so we do this in
uh kind of like a three-pronged approach
for directed instruction
so we give them a textbook reading that
they're
able to to use this is something we have
to update every year because the
graphical user interface that we choose
to use
changes we don't teach the command line

for this in the first year because it's
a bit too
overwhelming i think for the students um
then we do a live demonstration where
they
they watch us use the github website
they watch us use the get gui they watch
us move files and
add and commit and push and pull and
then finally they work through a guided
worksheet
that asks them to do the same thing that
we just did
and then ask them questions along the
way uh to test that they
like really understand like what is
committing what is adding what is
pulling what is pushing where is the
work going
and if folks are interested i have uh
some links embedded in this talk
that will take you to some of the
examples or resources that i'm talking
about here
okay and so the final thing that i said
um

Lot's of practice

is lots of practice so why do we need
lots of practice
for reproducibility for learning

reproducibility workflows and
tools well there are really two
fundamental ways that we commit things
to long-term memory
one is one trial learning and that
usually requires some sort of emotional
impact so that's like sometimes it's
traumatic events and sometimes it's
really good positive emotional events
that you had like a really
great birthday or your wedding or
something like that you don't need those
things to be repeated multiple times so
that you remember them
but that's not most of the things in
school most of the things in school
we learn about are through this
repetitive space training
um and so uh the the you know the best
way to commit
uh something to long-term memory that's
not really emotional
is to revisit it and repeat it multiple
times and have breaks between
those things and so that lets you commit
it to memory
however i think you want to go even a
step further with reproducibility
because when we teach reproducibility
workflows and skills as instructors
we actually want students to do more

than just learn about these
things we actually want them to use them
and put them into practice
um in the classroom outside of the
classroom in their work after the
classroom
and so we actually want to change their
habits or behaviors and it's it's quite
i think
important to realize that okay it's not
just understanding an algorithm
it's it's understanding the concepts
behind something like version control
understanding the concepts behind
something like a shippable and shareable
compute environment
and then knowing how to use those things
and then
once you leave the classroom wanting and
being able to use those things without
like
saying no that's too hard or too tricky
you want them to just
do it out of habit because that's what
they usually do
um so an aside uh just a little
if people are interested a book that's
really recently made me think about more
how we can tangibly do this is called
atomic habits by james clear
he's done a really good job of like

bringing the science of
habit building and behavior change uh
into an accessible book
and um i think that when we think about
getting students practice and changing
their behavior with
reproducible skills and workflows
there's a lot of really
interesting insight from behavior change
and psychology and habit building that
we can that we can borrow
so um now i'll talk more practically
about so
at least right now what are we doing in
the classroom to embed
lots of practice so
what we do is when we do our live demos
we don't just have
us do it then we pause and say okay
students your turn
do what i just did and so they saw
it and then they actually have to type
it into the keyboard or click it their
mouse around
the graphical user interface so they
they practice it that
way then we have lots of low stakes
assessments
with small or short problems so um
we've moved into a lot of flipped
classroom

um in in at ubc or at least in the data science so our introduction to data science course is a primarily flipped classroom so um we have uh literate code documents that have uh automated tests in them that the students are answering all kinds of questions about the data science content and then they're very short little pieces that are well guided but they get immediate feedback um and these things aren't worth very much and they do a lot of them so they do two of them a week in the data science course in the master of data science um program we've also started implementing this in some of our classrooms and the students really like this practice and it helps them really prepare for things like larger assessments like quizzes and and their their lab homework but it gives them lots of practice and then the other thing that we do is the learning technologies and platforms that we use are built and use authentic data science reproducibility tools and so i'll give an example of that now so in almost all of the master beta

science courses so i'm talking about 20 courses here

21 credit courses

so 21 month long courses we use version control particularly github as our course management system

so the homework instructions and assignments are distributed to the students as github repositories

and the only way that they can submit their homework is by putting their homework in that github repository

so they have to go through the cloning procedure or at least be able to somehow download this from the github uh website and then they have to be able to hopefully through things like pushing and committing send their work back to github but they would at least have to interface with the github

uh website to do this um

to try and uh incentivize um the actual actually using the get machinery to interact with github

we also put part of the marks of each of these assignments as to mechanics and so um

we need to see for example like three commits associated with every single assignment because

um we think you know we're trying to build these good habits and practices around like there's reasons why you use version control not just to submit your homework but

to active as a backup or in case you want to go back in time and change things

so by the end of this program um the students have version controlled their work in over 80 different repositories um so they have a lot of so they're very practiced and very used to it and they're basically you know you want them to be able to do it in their sleep um almost and so that when they leave the program and they go to work somewhere else it's just natural it's just one of their habits at this point that if they're going to work on a project it's going to go under version control

so we do this using tools and here i've listened there's many tools now which is pretty cool we're not the only program doing this

at all there's many tools now for using github

as a classroom learning management system um and so i've listed a couple of them here

folks are interested so

Conclusions

We think over the past five years from teaching in the Master of Data Science and our Introduction to Data Science course that key things for teaching reproducibility in the data science classroom are providing extra emphasis on motivation, providing direct instruction so it's not a mystery of how you get started and what you need to do, as well as lots and lots of practice so that we can not only teach them the material and the concepts but so that this actually changes their practices and their workflows and they will use it after leaving your classroom.

Q & A's from talk

(leaving here for now in case there is anything inspirational to add to the paper)

so i'm happy to take questions now here

um

or uh you can tweet to me on twitter and

i'm happy to answer there and again i've

posted the link for the slides

thanks very much tiffany amy did you

want to share the question period or

shall i um i

can i'm just gonna check the thread

i don't see anything right now

anyone have anything to start off with

i have some if there's none um otherwise

maybe mina has

comments i'd be i'd love just to have

mina and tiffany just like

tell us everything you know between the

two of you

was there something that you started

doing that you i mean obviously

it's a evolution right the these these

all of these
programs are just in the evolution stage
uh was this uh
something that you started doing that
you've really moved away from
yeah i would say the pandemics even
placed a greater emphasis on this
so when we started teaching the master
of data science program
um we had a small cohort and we were in
person which allowed us to provide a lot
of support
but as you scale these things um
having that intimate close support is
more difficult and so in the very first
year of the program
um in the mbs program we have like very
we we have this philosophy that they
should be able to
be somewhat experts of running stuff on
their laptop we do teach them some cloud
tools but
you should be able to install your
software stack and and be able to set up
your path and these sorts of things
and so um we have a pretty i'd say like
intensive list of like 20 things that
they need to install in the first week
they need to use git to submit their
homework in the first week
and it's a bit overwhelming and it's a

lot and so
um what we've kind of moved we've kind
of like eased off on that
and and moved um to we get them there
but we take long
we take longer now to doing that so uh
we've set up this year a jupiter hub to
have them work in a cloud-based setting
for the first week or two
and then after the first week or two is
when we transition them to their own
laptop
so that we give time for the no like the
the
expectation that setting up everybody's
system
um is going to encounter some bugs and
take some time and that's going to be
tricky
we've also for the first assignment we
no longer ask them in the first
assignment to
submit to github um that's assignment
two so that we have to give ourselves
like a week or two
to to get them up to speed for getting
github
um and so i think that uh that sort of
thing
has um has definitely changed and been
inspired with

so first i started working in the master
data science program and then i started
teaching undergraduates
and teaching undergrad graduates has
made me have to reframe things and think
about things differently
um and think about like how do i remove
barriers so that
people you know maybe people who who are
for whatever reason more sensitive to
not feeling like data sciences for
them i don't want them to drop my class
because they couldn't install something
um so i think yeah that's something
that's changed a lot
let's quickly get to meena's question
where she needs to go um
a question about have you seen any
changes in computing experiences of
students applying to your
ms program um are more students coming
in with familiarity with these tools or
not yet
yeah yeah that's a great question i do a
little survey
every year about like in in the first
class
i'm like what tools have you used before
and usually about half of the students
have used r
maybe three quarters of the students

have used python almost all of them have
seen jupiter
almost none of them still have seen get
in github um so it's really quite
amazing
that um
computer science programming
prerequisite uh so they have
it's they don't have to have had our pro
or python before
um but yeah it's still interesting that
even
though i'll comment in a second that we
are seeing people with more
technical or data science skills coming
to our program
it's still the reproducibility
experience with reproducibility related
tools are
are aside from jupiter like um not
as present as one might expect i am
seeing more and more
uh people having like in data science
applying to our program which is
something kind of new and interesting
for us to think about because our
program was really designed
not for somebody who is like a data
science undergraduate like somebody who
had
an undergraduate in a different

discipline and wants to then apply
you know data science to their
discipline so we're still thinking about
like how
how we're going to handle the change of
like there's going to be more and more
undergraduates coming in with this
expertise
yeah it's a super interesting problem um
john's asking
um what to do with docker and windows
um is there something special about
documents yeah
so it uh it can work
um but everything with windows is a
little bit more challenging
uh so what my strategy is is i have
um i i'm a mac person uh but i have a
i also have a separate pc laptop where i
have linux and windows installed
so that before i teach every course i go
through and make sure that
i know how to install things on windows
and
what instructions to provide to students
there's still always surprises
um one thing we do this on quite a large
scale largest scale with the master
students but i'd say like
we're dealing with 50 or 60 windows
different windows laptops every

every year and so to make our lives
easier
um we've been really tightly restricting
which version of windows
that they have because uh then it's
easier to know so we say you have to
have windows 10
you have to have this build um and it
can't be windows home um basically it
has to be enterprise pro
or or education and by doing that that
has reduced some problems but every year
something new comes up like
i can just tell you this week i'm
teaching uh python packaging with poetry
and git bash doesn't work with poetry
anymore this year it worked last year
but it doesn't work anymore there's a
game
that have issue open it's not resolved
so now we're using anaconda prompt on
the windows machines
we have a solution um but it's
it's just it's it's one of those there
there will be dragons in this field
yeah it's yeah keeping changing things
it's just so much work right
you think you're done and at least 20
years ago right the folks they write out
their theory equations and that was it
they were done for the next 20 years

we've got to update ourselves every six
months
oh did you have any other closing
comments or thoughts that you wanted to
say